

BlueOwl

February 4, 2020

Last name: ILORI Date: 04-February-2020 Approach:Upsampled the minority class since the dataset had imbalanced classes and then modeled on the upsampled dataset and predicted on the original test set. Target AUC: 0.825 Achieved AUC: 0.974

```
[1]: #import packages
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[2]: #load the datasets
train_data = pd.read_csv('train.csv')
test_data = pd.read_csv('test.csv')
```

```
[3]: #display first five observations of both data
train_data.head(5)
test_data.head(5)
```

```
[3]:   age  cost_of_ad  device_type  gender  in_initial_launch_location  income  \
0   34    0.005134    Android      F                                1    40376
1   53    0.005223    desktop      F                                1    84511
2   46    0.004939    laptop      F                                0    79322
3   36    0.004924    Android      F                                0    63295
4   28    0.005146     other      F                                1    36170
```

```
   n_drivers  n_vehicles  prior_ins_tenure
0           1           3                 7
1           1           1                11
2           1           1                 4
3           1           2                 0
4           1           3                 3
```

```
[4]: #display info of the data
train_data.info()
test_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 10 columns):
age                                10000 non-null int64
```

```

cost_of_ad          10000 non-null float64
device_type         10000 non-null object
gender              9731 non-null object
in_initial_launch_location 10000 non-null int64
income              10000 non-null int64
n_drivers           10000 non-null int64
n_vehicles           10000 non-null int64
prior_ins_tenure     10000 non-null int64
outcome             10000 non-null int64
dtypes: float64(1), int64(7), object(2)
memory usage: 781.3+ KB
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 9 columns):
age                  10000 non-null int64
cost_of_ad           10000 non-null float64
device_type          10000 non-null object
gender               9751 non-null object
in_initial_launch_location 10000 non-null int64
income               10000 non-null int64
n_drivers            10000 non-null int64
n_vehicles            10000 non-null int64
prior_ins_tenure      10000 non-null int64
dtypes: float64(1), int64(6), object(2)
memory usage: 703.2+ KB

```

```

[5]: #view description of the data
train_data.describe()
test_data.describe()

```

```

[5]:
      count  age  cost_of_ad  in_initial_launch_location  income \
count  10000.000000  10000.000000  10000.000000  10000.000000
mean    38.029200    0.004668    0.500900  58232.104800
std     12.896921    0.000836    0.500024  17318.192722
min     16.000000    0.003173    0.000000  16124.000000
25%     27.000000    0.004000    0.000000  45575.250000
50%     38.000000    0.004609    1.000000  58271.000000
75%     49.000000    0.005175    1.000000  70438.750000
max     60.000000    0.007049    1.000000  99916.000000

      count  n_drivers  n_vehicles  prior_ins_tenure
count  10000.000000  10000.000000  10000.000000
mean    1.486300    2.009100    5.319800
std     0.499837    0.817242    5.019487
min     1.000000    1.000000    0.000000
25%     1.000000    1.000000    1.000000
50%     1.000000    2.000000    4.000000
75%     2.000000    3.000000    8.000000

```


```
max          2.000000      3.000000      22.000000
```

```
[6]: #display distribution plot for the outcome feature
fig, ax = plt.subplots(figsize=(4, 4))
sns.distplot(train_data.outcome, ax=ax)
```

```
C:\Users\Israel Ilori\Anaconda3\lib\site-packages\scipy\stats\stats.py:1713:
FutureWarning: Using a non-tuple sequence for multidimensional indexing is
deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will
be interpreted as an array index, `arr[np.array(seq)]`, which will result either
in an error or a different result.
```


```
    return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

```
[6]: <matplotlib.axes._subplots.AxesSubplot at 0x21828161a20>
```




```
[7]: #using boxplot check for outliers
import seaborn as sns
sns.boxplot(x='outcome', data = train_data)
```

```
[7]: <matplotlib.axes._subplots.AxesSubplot at 0x2182a361a20>
```



```
[8]: #correlation between features in the data  
a = train_data.iloc[:, :9]  
corr_matrix = a.corr()  
plt.subplots(figsize=(10,6))  
sns.heatmap(corr_matrix, vmax=0.9, square=True)
```

```
[8]: <matplotlib.axes._subplots.AxesSubplot at 0x2182a3b3908>
```



```
[9]: #Feature encoding for categorical values in the dataset
from sklearn.preprocessing import LabelEncoder

number = LabelEncoder()
train_data['gender'] = number.fit_transform(train_data['gender'].astype('str'))
test_data['gender'] = number.fit_transform(test_data['gender'].astype('str'))
train_data['device_type'] = number.fit_transform(train_data['device_type'].
→astype('str'))
test_data['device_type'] = number.fit_transform(test_data['device_type'].
→astype('str'))
```

```
[10]: #display counts for both classes
train_data['outcome'].value_counts()
```

```
[10]: 0    9018
      1     982
      Name: outcome, dtype: int64
```

```
[11]: from sklearn.utils import resample

#Separate majority and minority classes
train_majority = train_data[train_data.outcome==0]
train_minority = train_data[train_data.outcome==1]

#assign variable count of the majority class to be used in the method
majority_count = len(train_data[train_data.outcome==0])

#Upsample minority class
train_minority_upsampled = resample(train_minority,
                                   replace=True,      #sample with replacement
                                   n_samples=majority_count,  #to match
                                   ↪majority class
                                   random_state=123) # reproducible results

#Combine majority class with upsampled minority class
train_upsampled = pd.concat([train_majority, train_minority_upsampled])

#Display new counts for both classes
train_upsampled.outcome.value_counts()
```

```
[11]: 1    9018
      0    9018
      Name: outcome, dtype: int64
```

```
[12]: from sklearn.linear_model import LogisticRegression
      from sklearn.metrics import accuracy_score, f1_score, auc, roc_curve,
      ↪roc_auc_score
      from sklearn.ensemble import RandomForestClassifier

#Separate input features (X) and target variable (y)
y = train_upsampled.outcome
X = train_upsampled.drop('outcome', axis=1)
test = test_data.iloc[:,:].values

from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split (X, y, test_size = 0.3,
      ↪random_state = 42)
```

```
[13]: #Train model
      #Instantiate model and fit
      log_reg = LogisticRegression().fit(X_train, y_train) #serves as base model
      rand_forest = RandomForestClassifier().fit(X_train, y_train) #serves as second
      ↪model

#Predict on train set
```

```
y_pred = rand_forest.predict(X_test)

#display the classes for which model is predicting
print(np.unique(y_pred))

#display performance metrics
print('The Accuracy score is: ',accuracy_score(y_test, y_pred))
print('The F1_score score is: ',f1_score(y_test, y_pred))
print('The AUC score is: ',roc_auc_score(y_test, y_pred))
```

```
[0 1]
The Accuracy score is:  0.9732027351690999
The F1_score score is:  0.9740932642487047
The AUC score is:  0.9729981378026071
```

```
[14]: #Run prediction on the test set
y_pred2 = rand_forest.predict(test)

#Display the first five predicted values
y_pred2[:5]
```

```
[14]: array([0, 0, 0, 0, 0], dtype=int64)
```

```
[ ]:
```